

Performance through caching

Aurélien Broszniewski
@AurBroszniewski



BARCELONA JAVA USERS GROUP



Aurélien Broszniowski

Lead Engineer at Terracotta/Software AG

Ehcache!

Initiated the Rainfall framework.
Testing JSR107 caches.

Extension to a generic stress and
performance testing framework.



What we'll cover

- **JCache (jsr107)**
- **Cache Aside**
- **Cache-Through**
- **Performance testing**
- **Testing jsr107 providers**
- **Statistics**
- **Beyond JCache... capacity**
- **Beyond JCache... topology**


Introductory poll

- Who knows nothing about caching?
- Who already uses caching in production?
- Who had caching related problems in production?
- Who knows about JSR 107?

What is a cache?

- Data structure holding a *temporary* copy of some data
- Trade off between *higher memory* usage for *reduced latency*
- Targets:
 - Data which is reused
 - Data which is expensive to compute or retrieve

Why does caching work?



- Pareto distribution : The long tail
- 80/20

Why does caching work?

$$S_{\text{latency}}(s) = \frac{1}{(1 - p) + \frac{p}{s}}$$

where

- S_{latency} is the theoretical speedup in latency of the execution of the whole task;
- s is the speedup in latency of the execution of the part of the task that benefits from the improvement of the resources of the system;
- p is the percentage of the execution time of the whole task concerning the part that benefits from the improvement of the resources of the system *before the improvement*.

- Amdahl's law

Possible usages

- CPU bound applications
 - Normal speedup through algorithm improvement or parallelisation
 - Cache can help by storing computation results
- I/O bound applications
 - Normal speedup through disk or network upgrades
 - Cache can help by storing data locally

107

JSR 107
JCache

JSR 107

- Java Community Process driven standard
- Specifies API and semantics for **temporary, in-memory caching** of Java objects, including object creation, shared access, spooling, invalidation, and consistency across JVM's

javax.cache API

CacheManager repository,

```
CachingProvider provider = Caching.getCachingProvider();
```

```
CacheManager cacheManager = provider.getCacheManager();
```

```
Cache<Long, String> myCache = cacheManager.getCache("myCache", Long.class, String.class);
```


Named Cache repository,
handles their lifecycle

ConcurrentMap<K, V>
sibling, major interaction


Cache patterns

Cache aside

Cache aside - miss



Cache aside - hit




Exercise

[https://github.com/
aurbroszniowski/jbcn2016](https://github.com/aurbroszniowski/jbcn2016)


Cache patterns

Cache through


Cache through - miss



Cache through - hit



Cache through - write



Exercise

[https://github.com/
aurbroszniowski/jbcn2016](https://github.com/aurbroszniowski/jbcn2016)

Performance testing

Pitfalls

- **Warm-up phase (JIT optimisation!)**
- **Run reasonable length of time**
- **Do not average results!**
- **Think about amortization**
- **Pay attention to variability**

Common actions

- **Defining a Scenario of operations**
- **Executing the Scenario**
- **Gathering statistics**
- **Reporting results**

Rainfall framework

```
Runner.setUp(  
    Scenario.scenario("load test")  
        .exec(new Operation() {  
            @Override  
            public void exec(...) throws TestException {  
  
                long start = getTimeInNs();  
                // This is what we measure  
                service.someLogic(id);  
                //  
                long end = getTimeInNs();  
                statisticsHolder.record("READ", (end - start), READ);  
            }  
        })  
    .warmup(during(45, TimeDivision.seconds))  
    .executed(during(1, TimeDivision.minutes))  
    .config(report(Results.class).log(text(), html()))  
    .start();
```


Exercise

[https://github.com/
aurbroszniowski/jbcn2016](https://github.com/aurbroszniowski/jbcn2016)

JSR 107 providers

JSR 107

- **Ehcache**
- **Hazelcast**
- **Infinispan**
- **Apache Ignite**
- **more...**

Exercise

[https://github.com/
aurbroszniowski/jbcn2016](https://github.com/aurbroszniowski/jbcn2016)

Statistics

Statistics

- **HIT or MISS ? 50%!**
- **Statistics MBean**

```
javax.cache:type=CacheStatistics,CacheManager=urn.X-ehcache.jsr107-default-config,Cache=yourcache
```

Exercise

[https://github.com/
aurbroszniowski/jbcn2016](https://github.com/aurbroszniowski/jbcn2016)

Beyond JCache

Beyond JCache

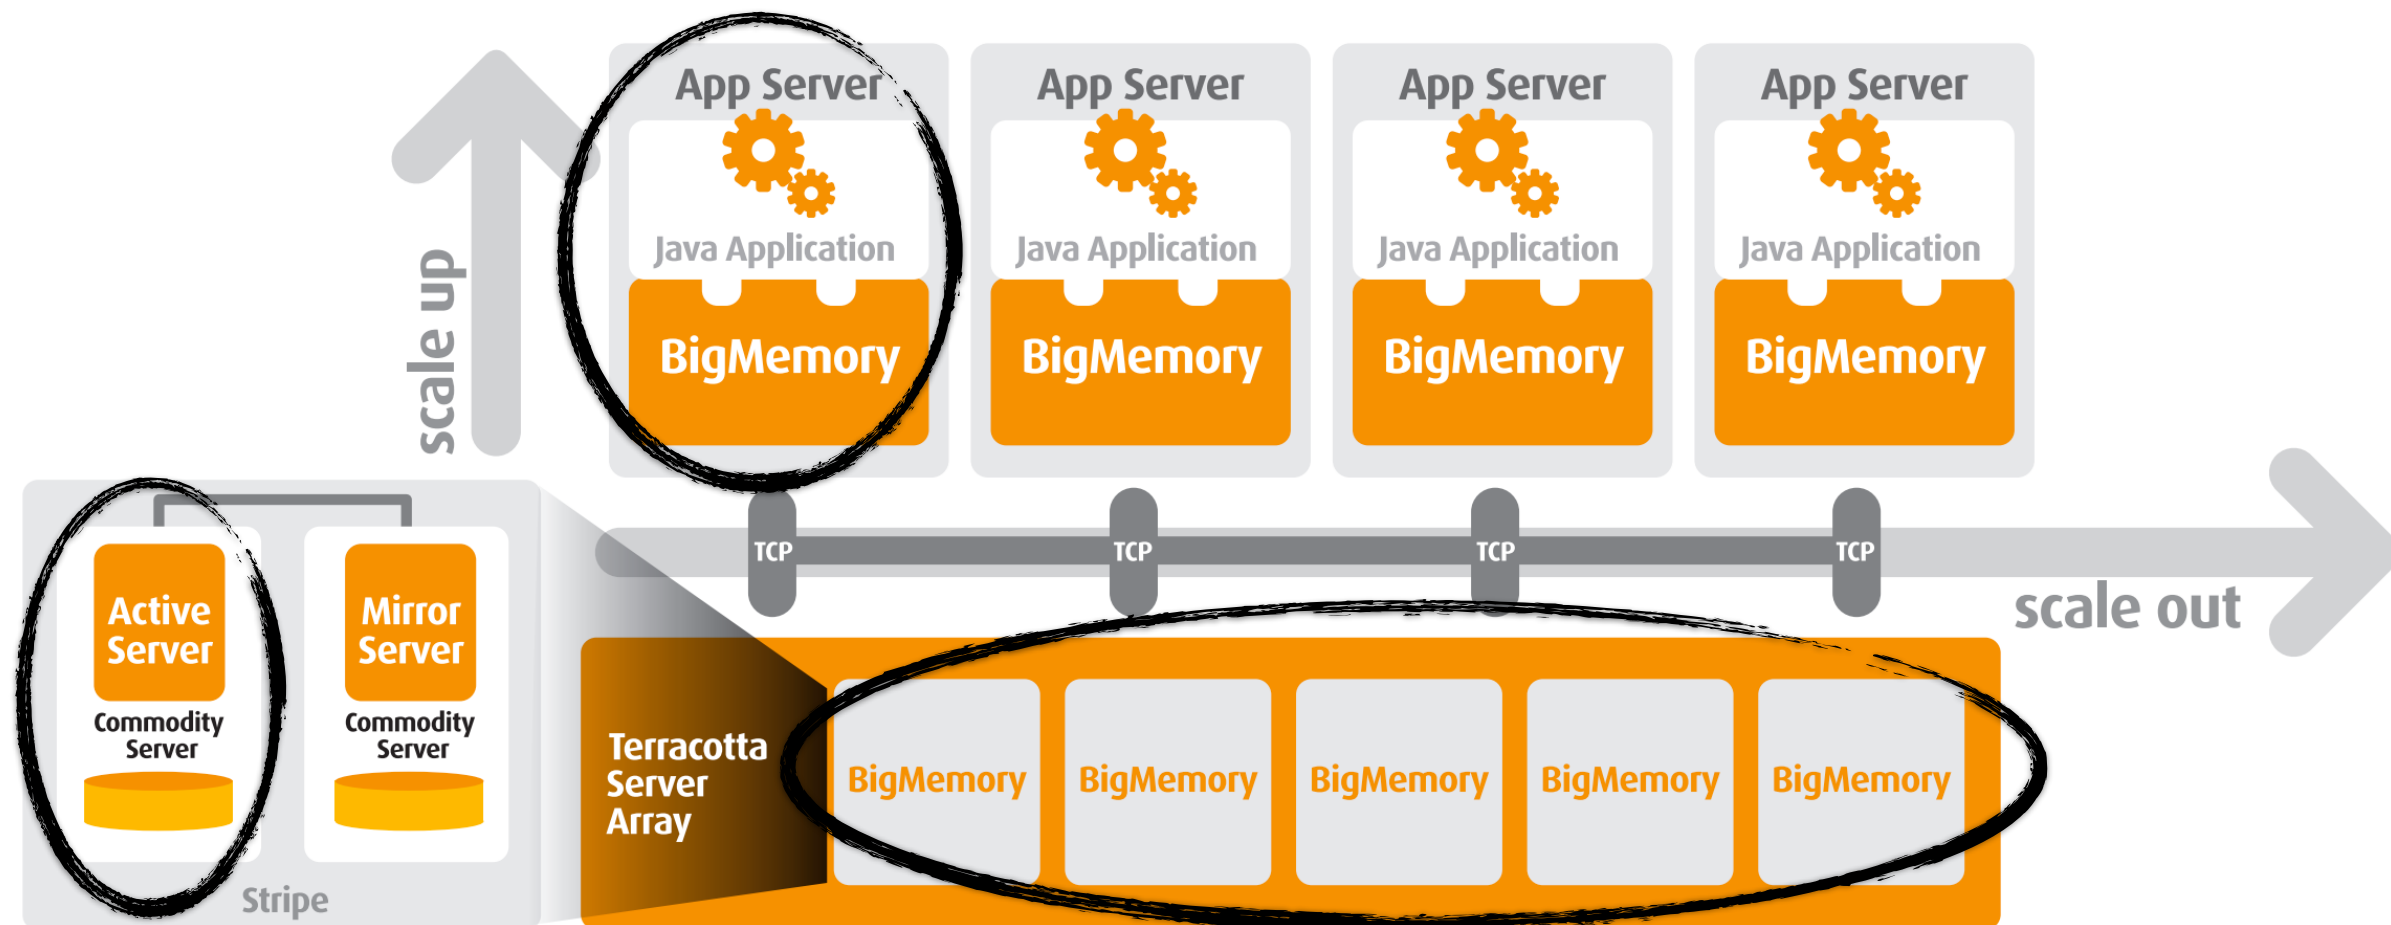
```
CacheConfiguration<String, String> cacheConfiguration =  
    CacheConfigurationBuilder.newCacheConfigurationBuilder(  
        String.class, String.class,  
        ResourcePoolsBuilder.heap(10000))  
    .build();  
  
cache = cacheManager.createCache("someCache",  
    Eh107Configuration.fromEhcacheCacheConfiguration(  
        cacheConfiguration));
```

Exercise

[https://github.com/
aurbroszniowski/jbcn2016](https://github.com/aurbroszniowski/jbcn2016)

Beyond JCache : Cache topologies

Terracotta clustering



www.ehcache.org

www.rainfall.io

Thanks!

@AurBroszniowski

Sitting through your hour long presentation today made me realize that my attention span is at most 35 minutes long.



your  cards
someecards.com